

# Machine Learning in Signal Processing

Winter Semester 2022/23

4. Linear Classification

07.11.2023

Prof. Dr. Vasileios Belagiannis

Chair of Multimedia Communications and Signal Processing

# Course Topics

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

1. Introduction.
2. Basics and terminology.
3. Linear regression.
4. **Linear classification.**
5. Performance evaluation.
6. Neural networks.
7. Deep neural networks.
8. Decision trees.
9. Ensemble models.
10. Random forests.
11. Clustering / Unsupervised learning.
12. Dimensionality reduction.
13. Support vector machines.
14. Recap and Q&A.
  - The exam will be written.
  - We will have an exam preparation test before the end of the year.

## Acknowledgements

Ideas and inspiration from:

- CSC311 Introduction to Machine Learning, University of Toronto.
- Introduction to Machine Learning: LMU Munich.
- Introduction to Machine Learning, CSAIL, MIT.
- CSE 574 Introduction to Machine Learning, University of Buffalo.
- Special thanks Arij Bouazizi, Julia Hornauer, Julian Wiederer, Adrian Holzbock and Youssef Dawoud for contributing to the lecture preparation.

# Last Lecture Recap

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Linear regression.
  - Normal equation.
  - Gradient descent.
- Convergence.
- Polynomial regression.
- Regularisation.
- Ridge regression.
- Lasso regression.

# Today's Agenda and Objectives

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Linear classifiers.
- Grouping.
- Decision Regions and Boundaries.
- Linear Classifiers.
  - Perceptron Learning Rule.
  - Logistic Regression.
  - Naive Bayes Classifier.
  - Linear Discriminant Analysis.

# Classification Task

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The goal of classification is to predict a class label  $y$  from a finite set of categories  $y \in \mathcal{Y} = \{1, \dots, n\}$  for a given input sample.  $n$  is the number of classes.
- For example, classifying images with blooms:



*Bloom (class 1)*



*No bloom (class 0)*

- We consider the following cases:
  - Binary classification, where  $\mathcal{Y} = \{0, 1\}$  or  $\mathcal{Y} = \{-1, +1\}$  and  $n = 2$ .
  - Multiclass classification, where  $\mathcal{Y} = \{1, \dots, n\}$  and  $n \geq 3$ .

# Classification Task (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Binary classification, where  $\mathcal{Y} = \{0, 1\}$  or  $\mathcal{Y} = \{-1, +1\}$  and  $n = 2$ .
  - Medical diagnosis: classification whether a patient has a specific disease based on medical scans (image/volume input).
  - Classification of e-mails into spam and non-spam (text input).
  - Anomalous sound detection (audio input).
- Multiclass classification, where  $\mathcal{Y} = \{1, \dots, n\}$  and  $n \geq 3$ .
  - Handwritten digit recognition (image input)
  - Classifying e-mails into multiple categories spam, work, family, insurance and banking (text input).
  - Voice recognition (audio input).
- We refer to the problem also as statistical classification.

# Grouping Classification Algorithms

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- A classification algorithm can be defined as a mapping function  $f: \mathcal{X} \rightarrow \mathbb{R}^n$  that assign continuous outputs to given data points  $x \in \mathcal{X}$ .
- *Why do we define the classifier with continuous output?*
- Continuous vs discrete output:
  - Continuous-valued cost functions are easier to optimize, e.g. gradient based optimization.
  - In general, the continuous output is more informative.
  - The continuous output can be transformed to discrete (class) values. The opposite does not work.

# Grouping Classification Algorithms (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- We differentiate between two types of classifier depending if the outputs are scores or probabilities.
  - Scoring classifiers.
  - Probabilistic classifiers.
- In addition, for every classifier we differentiate between two classification approaches:
  - Generative classifiers.
  - Discriminative classifiers.
  - The classification approach depends on whether the classifier learn the model that generated the data or only separates the specified classes given the input features.

# Scoring Classifiers

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- For binary classification ( $n = 2$ ), we have only one mapping function  $f(x) = f_{+1}(x) - f_{-1}(x)$  as discriminant / scoring function.
  - The class labels are  $\mathcal{Y} = \{-1, +1\}$ .
  - The discrete class value is chosen by  $g(x) = \text{sgn}(f(x))$  or by thresholding  $g(x) := [f(x) \geq \delta]$  with  $\delta = 0$ .
  - $|f(x)|$  is referred to as confidence.
- *How many scoring functions do we have for a multi-class classification problem with 5 class categories?*
- For  $n$ -class classification ( $n \geq 3$ ), we have  $n$  discriminant / scoring functions  $f_1, \dots, f_n: \mathcal{X} \rightarrow \mathbb{R}$ .
  - The class labels are  $\mathcal{Y} = \{1, \dots, n\}$ .
  - The scores  $f_1(x), \dots, f_n(x)$  are transformed into discrete class values by choosing the class with the maximum score:  $g(x) = \arg \max_{i \in \{1, \dots, n\}} f_i(x)$ .

# Probabilistic Classifiers

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Probabilistic classifiers are a special case of scoring classifiers. The mapping function can be denoted as  $f(x): \mathcal{X} \rightarrow [0, 1]^n$ .
- For binary classification ( $n = 2$ ), a single function  $f(x)$  is sufficient. The probability output is transformed to a discrete class value with  $g(x) := [f(x) \geq \delta]$  with  $\delta = 0.5$ .
- For  $n$ -class classification ( $n \geq 3$ ), we have  $n$  probability functions  $f_1, \dots, f_n: \mathcal{X} \rightarrow [0, 1]$  with  $\sum_i f_i = 1$ .
- The probabilities  $f_1(x), \dots, f_n(x)$  are transformed into discrete class values by choosing the class with the maximum probability:  $g(x) = \arg \max_{i \in \{1, \dots, n\}} f_i(x)$ .

# Generative Classifiers

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Generative models aim to model both hidden and observed variables, e.g. input / output variables.
- Generative models assume that the distribution underlying the data follows a specific parametric form → parametric probability density estimation.
  - Our continuous (random) variables have a probability distribution. The probability density function (PDF) models the relationship between the input / output variables.
  - The shape of PDF defines the probability distribution function, e.g. Gaussian distribution.
  - Probability density estimation: we rely on the values of the observed variables to estimate the density of the probabilities.
- *Name an example of a generative model relationship.*

Reference to Density Estimation: <https://machinelearningmastery.com/probability-density-estimation/>

---

# Generative Classifiers (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Generative classifiers model the conditional distribution  $p(x|y = i)$ . Therefore, generative classifiers apply the Bayes theorem:
  - $f_i(x) = \mathbb{P}(y = i|x) = \frac{\mathbb{P}(x|y=i)\mathbb{P}(y=i)}{\mathbb{P}(x)} = \frac{p(x|y = i)p(y=i)}{\sum_{j=1}^n p(x|y = j)p(y=j)}$
- $f_i$  are the class probabilities estimated from the training data.
- Posterior probability  $\mathbb{P}(y = i|x)$ : the conditional probability of  $y = i$  given that  $x$  occurs.
- Likelihood  $\mathbb{P}(x|y = i)$ : the conditional probability of  $x$  given that  $y = i$  occurs.
- Prior  $\mathbb{P}(y = i)$ : the prior probability of  $y = i$ .
- Evidence  $\mathbb{P}(x)$ : the probability of  $x$  occurring.

# Discriminative Classifiers

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- There is not any assumption on the underlying data distribution  
→ distribution free.
- They focus on a direct mapping between input/output variables, while it also learns the conditional probability (the posterior, as we just saw)  $\mathbb{P}(y = i|x)$ .
- Discriminative classifiers do not model the underlying probability distributions.
- Discriminative classifiers directly learn the decision boundaries between classes using empirical risk minimization.
  - They rely on empirical risk minimization with a defined loss function to directly learn the quantity of interest.

# Empirical Risk Minimization

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- In supervised learning, the cost function is normally defined as follows:
  - $\mathcal{J}(\mathbf{w}) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{data}} \mathcal{L}(f(\mathbf{x}; \mathbf{w}), y)$ , where  $\mathcal{L}$  is the loss function,
  - $f(\mathbf{x}; \mathbf{w})$  the prediction function,  $\mathbf{x}$  the input and  $y$  the label.
- $\hat{p}_{data}$  is the empirical data distribution, which corresponds to the training data.
- Assuming access to the data generating distribution  $p_{data}$ , the cost function becomes:
  - $\mathcal{J}^*(\mathbf{w}) = \mathbb{E}_{(\mathbf{x}, y) \sim p_{data}} \mathcal{L}(f(\mathbf{x}; \mathbf{w}), y)$ .
- This is an ideal cost function that is not realistic due to the lack of access to  $p_{data}$ . It measures the expected generalization error that is usually called risk.

# Empirical Risk Minimization (Cont.)

---

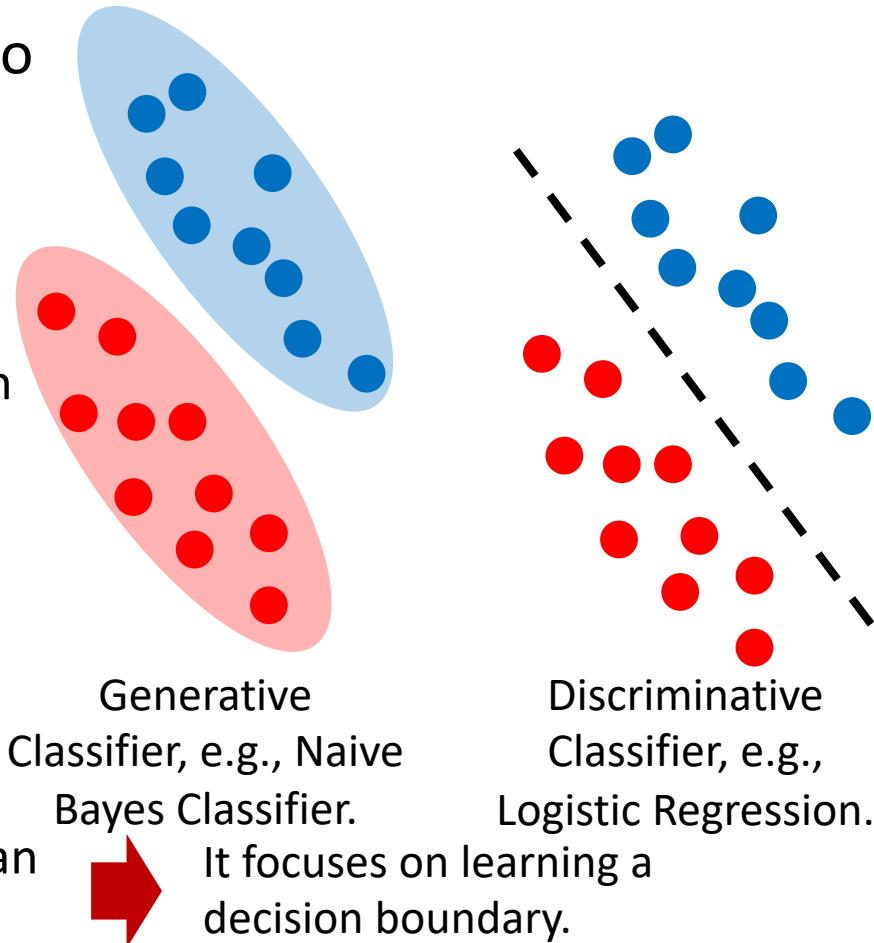
\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Our goal is to minimise the risk  $\mathcal{J}^*(\mathbf{w})$  but this is not possible.
- Instead we minimize the empirical risk as given by the cost function  $\mathcal{J}(\mathbf{w})$ . Given the training set  $\mathcal{T} = \{\mathbf{x}_i, y_i\}_{i=1}^m$ , we write the empirical risk as:
  - $\mathcal{J}(\mathbf{w}) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{data}} \mathcal{L}(f(\mathbf{x}; \mathbf{w}), y) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(f(\mathbf{x}_i; \mathbf{w}), y_i)$ .
- Training the discriminative classifier or some other machine learning algorithm results in minimizing the average error in the training set.
- In total, we minimize the empirical risk and hope that the risk is minimized too.
- What is a major disadvantage of the empirical risk minimization (ERM)?
- A disadvantage for the ERM is that it can overfit to the training data.

# Generative & Discriminative Classifiers

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The discussed ideas generalise to different machine learning algorithms.
- Example:
  - Our task is to recognise the English and German text.
  - The generative classifier will learn the English and German language first and then decide for the language.
  - The discriminative classifier will learn separate English from German without learning the language.



Comparison between the two models: <https://www.baeldung.com/cs/ml-generative-vs-discriminative>

# Decision Regions and Boundaries

---

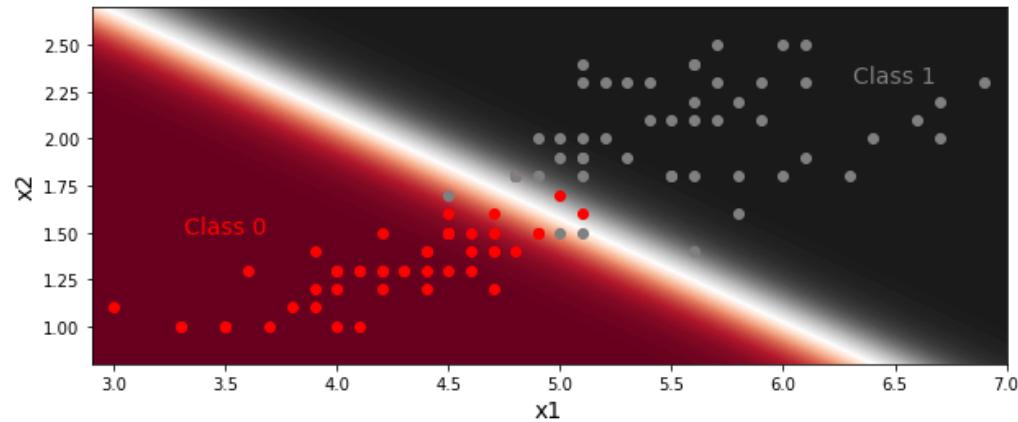
\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Decision region: set of data points  $\mathbf{x}$  to which the model assigns the class  $y$ :  $\mathcal{X}_y = \{\mathbf{x} \in \mathcal{X} : g(\mathbf{x}) = y\}$  where  $g(\cdot)$  the mapping function.
- The decision boundary separates the points of two or more classes.
- In the binary case, the decision boundary is given by:  $\{\mathbf{x} \in \mathcal{X} : f(\mathbf{x}) = c\}$ . It is a threshold.
- In the multiclass case, the decision boundary is defined as:  $\{\mathbf{x} \in \mathcal{X} : \exists i \neq j \text{ s. t. } f_i(\mathbf{x}) = f_j(\mathbf{x}) \text{ and } f_i(\mathbf{x}), f_j(\mathbf{x}) \geq f_k(\mathbf{x}) \forall k \neq i, j\}$ .

# Decision Regions and Boundaries (Cont.)

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The figure illustrates the hypersurfaces of two examples classes “0” and “1” with corresponding decision boundary as defined by a linear classifier.
- In the case of two classes, we have a linear binary classifier.

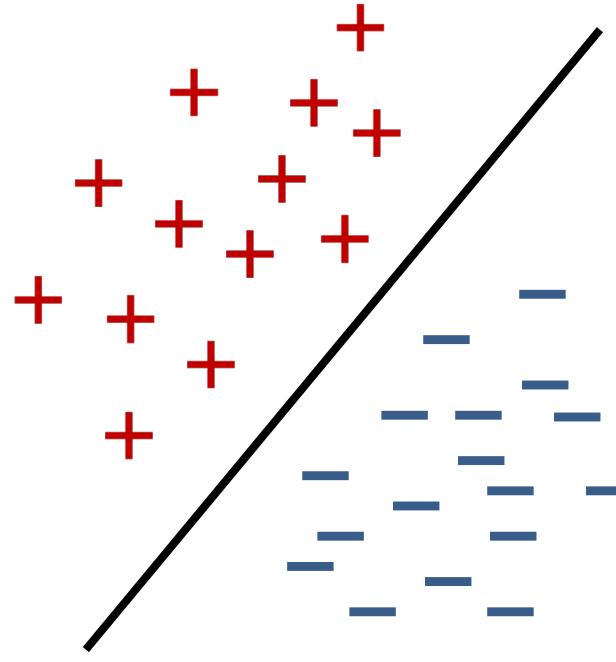


# Linear and Non-Linear Decision Boundary

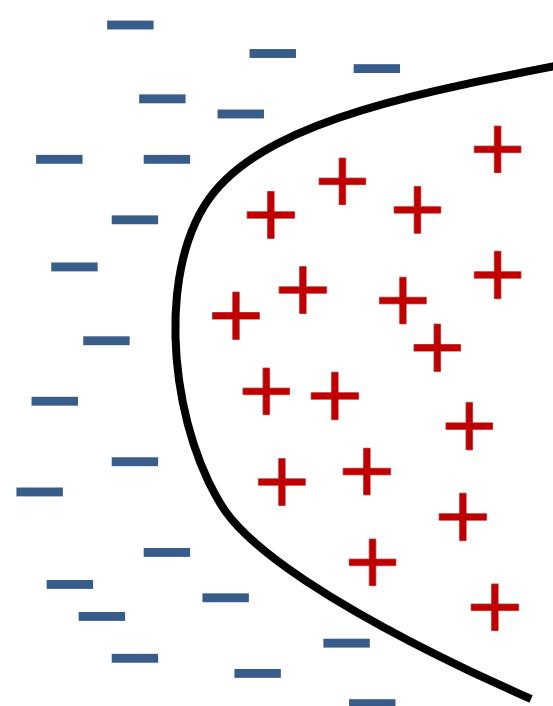
\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Binary Classifier.

Linear Decision Boundary



Non-Linear Decision Boundary



# Linear Classifiers

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- A classifier is linear if the mapping function(s)  $f_i(\mathbf{x})$  can be either defined directly as linear function(s) or through a rank-preserving monotone transformation  $g: \mathbb{R} \rightarrow \mathbb{R}$ .
  - $f_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x} + \mathbf{b}_i$ .
  - $g(f_i(\mathbf{x})) = \mathbf{w}_i^T \mathbf{x} + \mathbf{b}_i$ .
- Rank-preserving implies that the decision region and the decision boundary do not change after applying the transformation  $g$ .
- For linear classifiers, the decision boundary between two classes  $i$  and  $j$  is a hyperplane separating the two classes:
  - $f_i(\mathbf{x}) = f_j(\mathbf{x})$
  - $g(f_i(\mathbf{x})) = g(f_j(\mathbf{x}))$
  - $\mathbf{w}_i^T \mathbf{x} + \mathbf{b}_i = \mathbf{w}_j^T \mathbf{x} + \mathbf{b}_j$
  - $(\mathbf{w}_i - \mathbf{w}_j)^T \mathbf{x} + (\mathbf{b}_i - \mathbf{b}_j) = 0$
- This hyperplane is  $(\mathbf{w}_i - \mathbf{w}_j)^T \mathbf{x} + (\mathbf{b}_i - \mathbf{b}_j)$ .

# Linear Classifiers: Threshold Logic Unit

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Binary input  $x_i \in \{0,1\}$  and binary output  $y \in \{0,1\}$ .
- Excitatory or inhibitory input, represented by the not trainable weights  $w_i \in \{-1,1\}$ .
- The mapping is represented by a threshold function  $f: \mathbb{R}^d \rightarrow \mathbb{R}$ , where:
  - $f(\mathbf{x}) = \begin{cases} 0, & \text{if } \mathbf{w} \cdot \mathbf{x} \leq T \\ 1, & \text{otherwise.} \end{cases}$
- This is a linear binary classifier that was used to model gate functions.

---

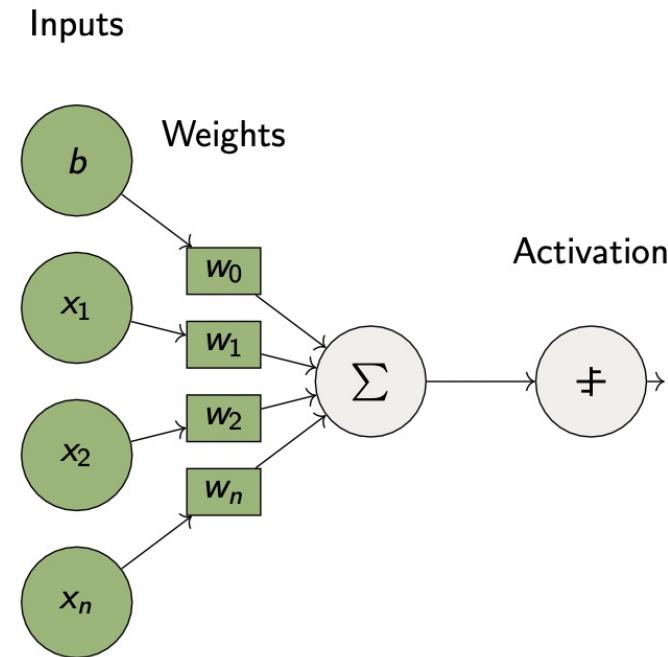
McCulloch, Warren S., and Walter Pitts. "A logical calculus of the ideas immanent in nervous activity." *The bulletin of mathematical biophysics* 5.4 (1943): 115-133.

# Linear Classifiers: Perceptron Learning Rule

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- It is a linear classifier for binary problems.
- We have a real-valued input and trainable weights / parameters  $w_i$ .
- Training data is the ground-truth is required  $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^m$ .
- It is represented by a threshold function  $f: \mathbb{R}^d \rightarrow \mathbb{R}$ , where:

$$f(\mathbf{x}) = \begin{cases} 1, & \text{if } \mathbf{w} \cdot \mathbf{x} + b \geq 0 \\ 0, & \text{otherwise.} \end{cases}$$



Rosenblatt, Frank. "The perceptron: a probabilistic model for information storage and organization in the brain." *Psychological review* 65.6 (1958): 386.

# Linear Classifiers: Perceptron Learning Rule (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Learning rule: the parameter  $w$  update is proportional to the input; and the difference between prediction  $f(\mathbf{x})$  and ground-truth  $y_i$ .
- Algorithm for single-layer Perceptron:
  1. Initialize the parameters  $w$ , set learning rate  $\eta$  and threshold.
  2. Update the weights:  $w_j(t + 1) = w_j(t) + \eta(y - f(\mathbf{x}))\mathbf{x}$
  3. Iterate the update until convergence.

Online example: <https://vitalflux.com/perceptron-explained-using-python-example/>

---

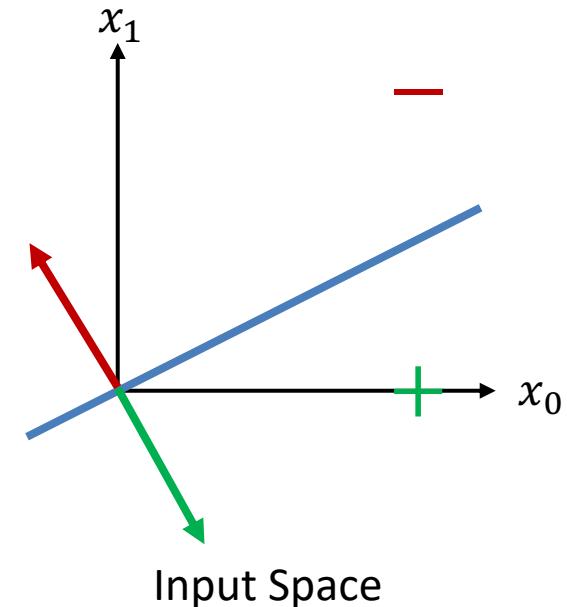
# NOT Gate

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- For the NOT gate only a single input  $x_1$  is required.
- The weights / parameters can be represented by *half-spaces*:  $H_+ = \{\mathbf{x}: \mathbf{w}^T \mathbf{x} \geq 0\}$  and  $H_- = \{\mathbf{x}: \mathbf{w}^T \mathbf{x} < 0\}$ .
  - Half-space: the set of data points from the one side of the hyper-plane.
- The boundary  $\mathbf{x}: \mathbf{w}^T \mathbf{x} = 0$  is the decision boundary.
- The training data is linearly separable, since we can find a decision boundary that correctly classifies all data points.
- An auxiliary input  $x_0$  is added for the geometric visualization.

$x_1$	$y$
0	1
1	0

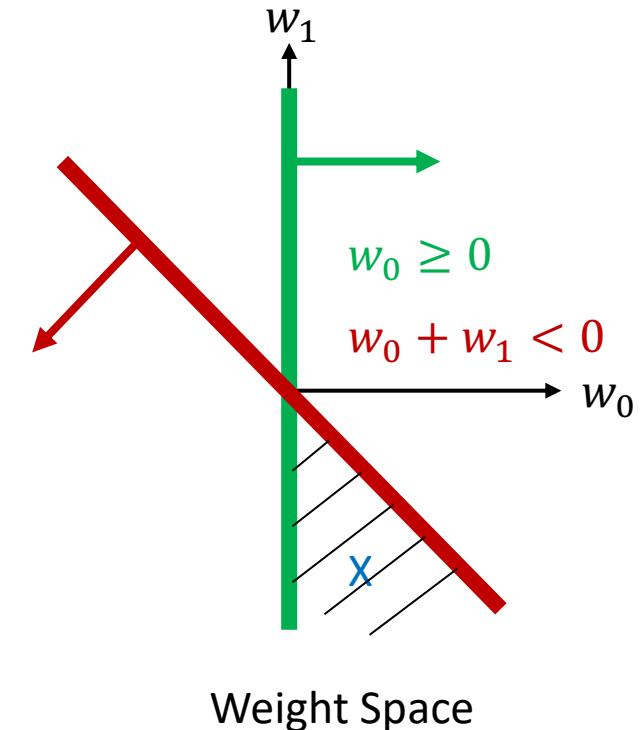
Training Set



# NOT Gate (Cont.)

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

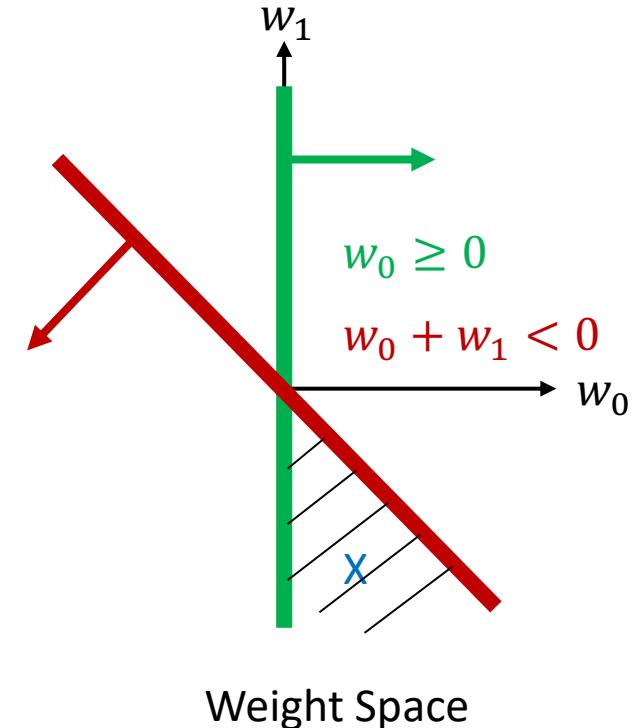
- The weights should fulfill the following conditions:
  - For  $x_1 = 0$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 \geq 0 \leftrightarrow w_0 \geq 0$
  - For  $x_1 = 1$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 < 0 \leftrightarrow w_0 + w_1 < 0$
- Solution:  $w_0 = 1$  and  $w_1 = -2$ .
- Each training sample  $\mathbf{x}$  specifies a half-space  $\mathbf{w}$ . The training example should lie within the half-space to be correctly classified.



# NOT Gate (Cont.)

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- In our example:
  - For  $x_0 = 1, x_1 = 0, y = 1$  we have  $(w_0, w_1) \in \{w: w_0 \geq 0\}$
  - For  $x_0 = 1, x_1 = 1, y = 0$  we have  $(w_0, w_1) \in \{w: w_0 + w_1 < 0\}$
- The region satisfying all above constraints is the feasible region.
  - The weights, which classify correct all the training samples, correspond to the intersection of all half-spaces.



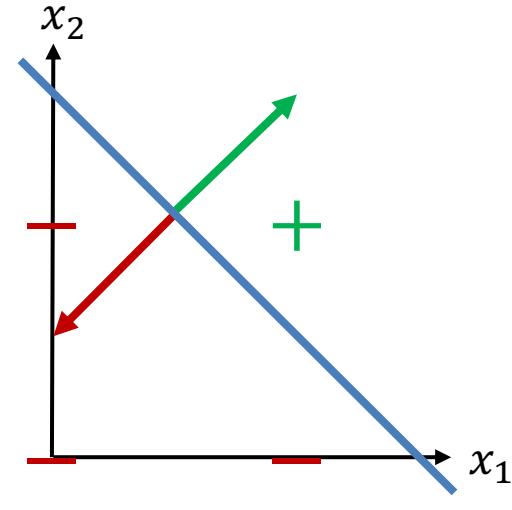
# AND Gate

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The AND gate requires two inputs  $x_1$  and  $x_2$ .
- To find a solution, a third auxiliary dimension is required.
  - The auxiliary input is always one.

$x_1$	$x_2$	$y$
0	0	0
0	1	0
1	0	0
1	1	1

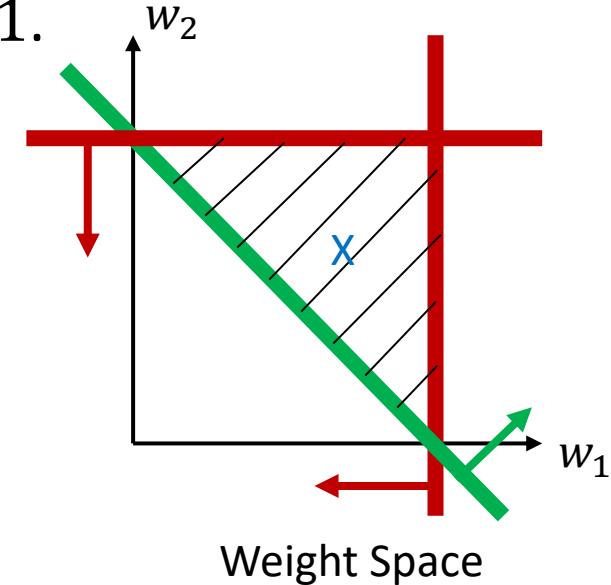
Training Set



# AND Gate (Cont.)

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Conditions for selecting the weights:
  - For  $x_1 = 0, x_2 = 0$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 + w_2x_2 < 0 \Leftrightarrow w_0 < 0$
  - For  $x_1 = 0, x_2 = 1$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 + w_2x_2 < 0 \Leftrightarrow w_0 + w_2 < 0$
  - For  $x_1 = 1, x_2 = 0$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 + w_2x_2 < 0 \Leftrightarrow w_0 + w_1 < 0$
  - For  $x_1 = 1, x_2 = 1$ :  $f(\mathbf{x}) = w_0x_0 + w_1x_1 + w_2x_2 \geq 0 \Leftrightarrow w_0 + w_1 + w_2 \geq 0$
- Solution:  $w_0 = -1.5$ ,  $w_1 = 1$  and  $w_2 = 1$ .
- The auxiliary dimension  $x_0$  is necessary for finding a solution.
- *Is there only a single solution?*



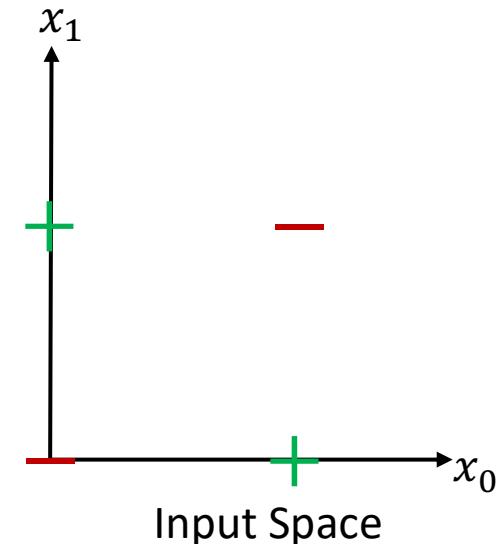
# XOR Gate

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Like AND, XOR has two input dimensions  $x_1$  and  $x_2$ .
- *Can we rely on a linear classifier for modeling the XOR gate?*

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0

Training Set



# Logistic Regression

---

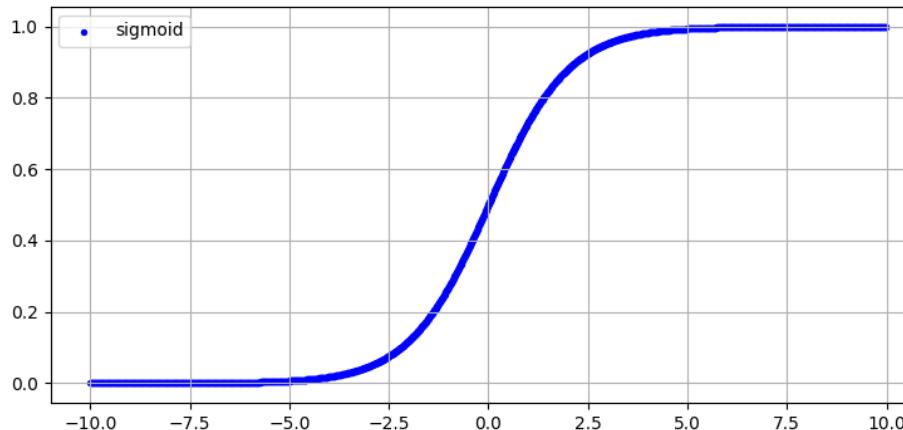
\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Logistic regression transforms the problem of linear regression to classification with the support of a function with bounded output.
- A few examples are spam mail detection, tumor classification in malignant or benign and online transaction classification to fraud or not.
- Motivation: convert continuous output of a bounded function to classification (i.e., category prediction).

# Logistic Regression (Two-Class Problem)

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Consider a two-class problem:  $y \in \{0, 1\}$ .
- We would like to map our input  $\mathbf{x}$  to one of these two values.
- Consider also the logistic function  $\sigma(\cdot)$ .



- The horizontal axis represents the input and the vertical axis the function's response.
- The value at 0.5 would be a well-suited threshold to decide between two classes, i.e., 0 and 1.

# Logistic Regression (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Recall the logistic function (also known as sigmoid) equation and its derivative:

- $\sigma(z) = \frac{1}{1+e^{-z}}$
  - $\frac{\partial \sigma}{\partial z} = \sigma(z)(1 - \sigma(z))$ , where  $z$  is a scalar.

- We deal with a regression problem, and we reformulate  $\sigma(\cdot)$  to be parameterized by  $\mathbf{w}$ , i.e., our parameters. We define the modified logistic function now as:

- $h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1+e^{-\mathbf{w}\mathbf{x}}}$

- The input  $\mathbf{x}$  is the feature vector.
- In addition, we interpret the output of the modified logistic function  $h_{\mathbf{w}}(\cdot)$  as the probability of being class 1. The function output does not have to be exactly 1 to classify the input  $\mathbf{x}$  as class 1.
- We can set a threshold at 0.5.

# Logistic Regression (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The probabilistic interpretation of the logistic function is written as:
  - $h_w(x) = p(y = 1|x; w) = 1 - p(y = 0|x; w)$ .
- We could rely on a set of  $m$  training pairs (input, ground-truth) and minimize the mean-squared error, like linear regression, where:

$$MSE = \frac{1}{m} \sum_{i=1}^m (h_w(x) - y_i)^2.$$

- Unfortunately, there is no **analytical solution** as with linear regression.
- We face a non-convex problem because of the sigmoid function.
- Instead, we can use the maximum likelihood estimation.

# Logistic Regression (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- We want to maximize the conditional probability function  $p(y|\mathbf{x}; \mathbf{w})$ .
- The conditional maximum likelihood estimation is given by:
  - $\mathbf{w}^* = \arg \max_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m \log p(y_i|\mathbf{x}_i; \mathbf{w})$ .
- The probability of  $\mathbf{x}_i$  to belong in class 0 or 1,  $y_i \in \{0, 1\}$ , given by:
  - $p(y_i = 1|\mathbf{x}_i) = h_{\mathbf{w}}(\mathbf{x}_i)$
  - $p(y_i = 0|\mathbf{x}_i) = 1 - h_{\mathbf{w}}(\mathbf{x}_i)$
- We can now rewrite the conditional maximum likelihood estimator:
  - $\mathbf{w}^* = \arg \max_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m y_i \log(h_{\mathbf{w}}(\mathbf{x}_i)) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i))$ .

# Logistic Regression (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Instead, we minimize the negative log-likelihood (NLL). This is written as:
  - $\mathbf{w}^* = \arg \min_{\mathbf{w}} -\frac{1}{m} \sum_{i=1}^m y_i \log(h_{\mathbf{w}}(\mathbf{x}_i)) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i))$ .
- This is a gradient-based optimization where we can use stochastic gradient descent.
- We assume the loss function:
  - $L(\mathbf{x}, y, \mathbf{w}) = \frac{1}{m} \sum_{i=1}^m -y_i \log(h_{\mathbf{w}}(\mathbf{x}_i)) - (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i))$ .
- The gradient w.r.t. the parameters is given by:
  - $\nabla_{\mathbf{w}} L(\mathbf{x}, y, \mathbf{w}) = \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i) \mathbf{x}_i$ .
- Logistic regression applies on the same way to multiclass problems as well.

Online example: <https://towardsdatascience.com/building-a-logistic-regression-in-python-step-by-step-becd4d56c9c8>

---

# Bayes Optimal Classifier

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- It is a probabilistic model that performs the most probable prediction for new sample, given the training data.
- $\mathbb{P}(y = i|x) = \frac{\mathbb{P}(x|y=i)\mathbb{P}(y=i)}{\mathbb{P}(x)} = \frac{p(x|y = i)p(y=i)}{\sum_{j=1}^n p(x|y = j)p(y=j)}.$
- In classification, we are interested in the class prediction. We aim for:
  - $\arg \max_{k \in \mathcal{Y}} \mathbb{P}(y = k|x).$
  - We can thus drop the evidence  $\mathbb{P}(x)$  and the posterior solution is still the same.
  - In general, we want find the  $x$  estimates that maximise the the posterior. This is known as Maximum a posteriori estimation (MAP).

Information on MAP: [https://www.probabilitycourse.com/chapter9/9\\_1\\_2\\_MAP\\_estimation.php](https://www.probabilitycourse.com/chapter9/9_1_2_MAP_estimation.php)

---

# Bayes Optimal Classifier

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Given a data point represent as feature vector  $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$  and label space  $\mathcal{Y} = \{0,1\}$ , we assume that  $x_i$  is in the label space.
- The optimal Bayes classifier is defined as:
  - $f_{Bayes}(\mathbf{x}) = \arg \max_{k \in \mathcal{Y}} \mathbb{P}(y = k | \mathbf{x})$
- The number of parameters increases with the number of features  $d$ , since  $2^d$  parameters are needed to describe the probability function  $p(y = k | \mathbf{x})$ , for example, when  $\mathbf{x} \in \{0,1\}^d$ .
- With the Bayes' theorem we only need to compute the posterior distribution  $p(\mathbf{x}|y = k)$ :
  - $f(\mathbf{x}) = \mathbb{P}(y = k | \mathbf{x}) = \frac{\mathbb{P}(\mathbf{x}|y = k)\mathbb{P}(y=k)}{\mathbb{P}(\mathbf{x})} = \frac{p(\mathbf{x}|y = k)p(y=k)}{\sum_{j=1}^n p(\mathbf{x}|y = j)p(y=j)}.$
- The optimal Bayes classifier can be written as:
  - $f_{Bayes}(\mathbf{x}) = \arg \max_{k \in \mathcal{Y}} p(\mathbf{x}|y = k)p(y = k).$

# Naive Bayes Classifier

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- The Naive Bayes classifier assumes that the features are conditionally independent given the class:
  - $p(\mathbf{x}|y = k) = p((x_1, x_2, \dots, x_d)|y = k) = \prod_{j=1}^d p(x_j|y = k).$
- Based on this assumption the Optimal Bayes classifier can be reformulated as:
  - $f_{Bayes}(x) = p(y = k) \prod_{j=1}^d p(x_j|y = k).$
- Then, only  $2d + 1$  parameters must be estimated.
- The required parameters are significantly reduced.

Vikramkumar, Vijaykumar, B., & Trilochan (2014). Bayes and Naive Bayes Classifier. *ArXiv*, *abs/1404.0933*.

---

# Naive Bayes Classifier (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Numerical feature:
  - The distribution  $p(x_j|y = k)$  is modelled as univariate Gaussian distribution.
  - Then, only the parameters  $\mu_j$  and  $\sigma_j^2$  must be estimated.
- Categorical features:
  - The distribution  $p(x_j|y = k) = \prod_g p_{kjg}^{[x_j=g]}$  is modelled as categorical distribution.
  - $p_{kjg}$  is the probability that the  $j$ -th feature has the value  $g$  in class  $k$ .
  - Therefore, the frequency of  $x_j = g$  can be counted.

Online example: <https://www.datacamp.com/tutorial/naive-bayes-scikit-learn>

---

# Naive Bayes Classifier (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- E-mail spam detector example.
  - Single words can be counted as feature to detect spam mails.
  - The classifier can be trained for each user individually. The probabilities assigned to each word are different for each user.
  - Independence assumption: the occurrence of two words in an e-mail is not correlated.
  - For example, the words “free money” occur with a high probability in spam e-mails, but the single words “free” and “money” also occur in non-spam e-mails with a high probability.
  - Individual letters of words that are frequently used in spam e-mails can be replaced with similar letters so that they are readable by humans but not recognized by spam filters.

# Linear Discriminant Analysis

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Again, a data point represent as feature vector  $\mathbf{x} = (x_1, x_2, \dots, x_d) \in \mathbb{R}^d$  and label space  $\mathcal{Y} = \{0,1\}$  is given.
- We now make the following assumptions:
  - $p(y = 1) = p(y = 0) = \frac{1}{2}$ .
  - $p(\mathbf{x}|y = k)$  is a Gaussian distribution, where the covariance matrix is the same for both class labels, i.e.  $\Sigma_k = \Sigma \forall k$ .
  - $p(\mathbf{x}|y = k) = \frac{1}{(2\pi)^{d/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_k)^T \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}_k)\right)$  with  $\boldsymbol{\mu}_0, \boldsymbol{\mu}_1 \in \mathbb{R}^d$  and  $\Sigma$  is the covariance matrix.

Tharwat, A., Gaber, T., Ibrahim, A., & Hassanien, A.E. (2017). Linear discriminant analysis: A detailed tutorial. *AI Commun.*, 30, 169-190.  
Online reference: <https://people.revoledu.com/kardi/tutorial/LDA/Numerical%20Example.html>

---

# Linear Discriminant Analysis (Cont.)

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Remember with Bayes rule the Bayes Optimal Classifier can be written as:  $f_{Bayes}(\mathbf{x}) = \arg \max_{k \in \mathcal{Y}} p(\mathbf{x}|y = k)p(y = k)$ .
- In the binary case,  $f_{Bayes}(x) = 1$  under the following condition:
  - $\log \left( \frac{p(x|y = +1)p(y=1)}{p(x|y = -1)p(y=0)} \right) > 0$ , which is called the *log-likelihood ratio*.
- Under the made assumptions the log-likelihood ratio becomes:
  - $\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_0)^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_0) - \frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_1)$ .
- This can be reformulated as a linear classifier:
  - $f(x) = \mathbf{w}^T \mathbf{x} + b$  with  $\mathbf{w} = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0)^T \boldsymbol{\Sigma}^{-1}$  and  $b = \frac{1}{2}(\boldsymbol{\mu}_0^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_0 - \boldsymbol{\mu}_1^T \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_1)$ .
- Now only the parameters  $\boldsymbol{\mu}_0$ ,  $\boldsymbol{\mu}_1$  and  $\boldsymbol{\Sigma}$  must be estimated from the training data.

Online example: <https://www.geeksforgeeks.org/ml-linear-discriminant-analysis/>

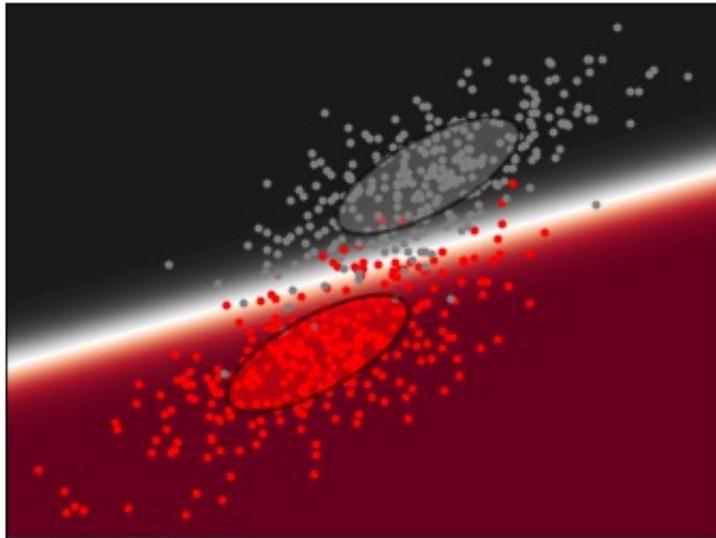
---

# Linear Discriminant Analysis (Cont.)

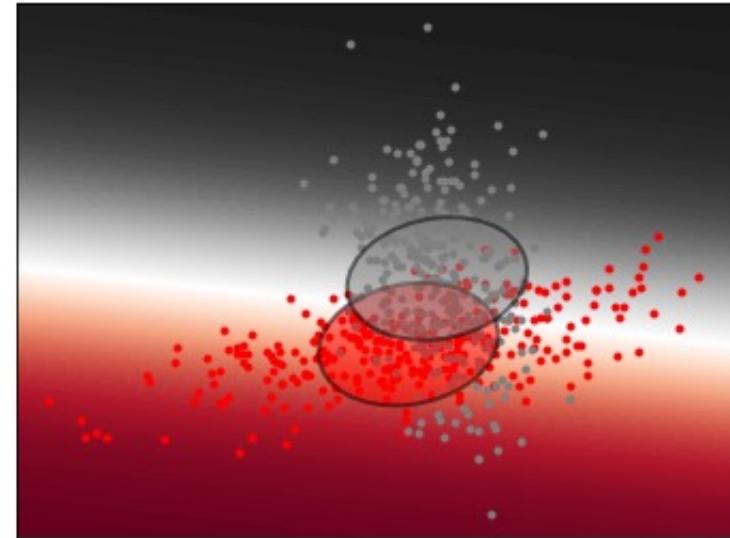
\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- Covariance matrix comparison

Same covariance matrix for both classes.



Different covariance matrix for both classes.



# Study Material

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

- *The Elements of Statistical Learning, Trevor Hastie et. al, Chapter 3.*
- *Understanding Machine Learning, Shai Shalev-Schwarz, Chapter 9, Section 9.2.*
- *Pattern Recognition and Machine Learning, Christopher Bishop, Chapter 3.*

# Next Lecture

---

\*\*\*Not for sharing (LMS, Friedrich-Alexander-Universität Erlangen-Nürnberg)\*\*\*

## Performance Evaluation